

Deep Neural Network Approaches to Speaker and Language Recognition

Fred Richardson, *Senior Member, IEEE*, Douglas Reynolds, *Fellow, IEEE*, and Najim Dehak, *Member, IEEE*

Abstract—The impressive gains in performance obtained using deep neural networks (DNNs) for automatic speech recognition (ASR) have motivated the application of DNNs to other speech technologies such as speaker recognition (SR) and language recognition (LR). Prior work has shown performance gains for separate SR and LR tasks using DNNs for direct classification or for feature extraction. In this work we present the application of single DNN for both SR and LR using the 2013 Domain Adaptation Challenge speaker recognition (DAC13) and the NIST 2011 language recognition evaluation (LRE11) benchmarks. Using a single DNN trained for ASR on Switchboard data we demonstrate large gains on performance in both benchmarks: a 55% reduction in EER for the DAC13 out-of-domain condition and a 48% reduction in C_{avg} on the LRE11 30 s test condition. It is also shown that further gains are possible using score or feature fusion leading to the possibility of a single i-vector extractor producing state-of-the-art SR and LR performance

Index Terms—Bottleneck features, DNN, i-vector, language recognition, senone posteriors, speaker recognition, tandem features.

I. INTRODUCTION

THE impressive gains in performance obtained using deep neural networks (DNNs) for automatic speech recognition (ASR) [1] have motivated the application of DNNs to other speech technologies such as speaker recognition (SR) and language recognition (LR) [2]–[11]. Two general methods of applying DNN's to the SR and LR tasks have been shown to be effective. The first or “direct” method uses a DNN trained as a classifier for the intended recognition task directly to discriminate between speakers for SR [5], [11] or languages for LR [4]. The second or “indirect” method uses a DNN possibly trained for a different purpose to extract data that is then used to train a secondary classifier for the intended recognition task. Applications of the indirect method have used a DNN to extract frame-level features [2], [3], [12], accumulate a multinomial or

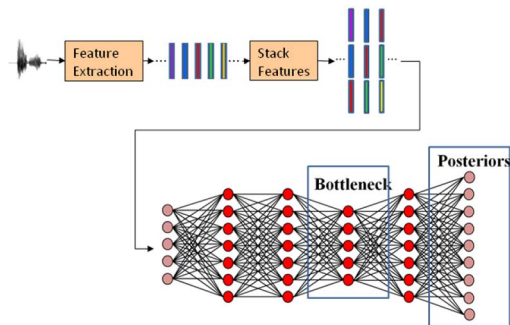


Fig. 1. Example DNN architecture.

Gaussian vector [7], [11] or accumulate multi-modal statistics [6], [8] that were then used to train an i-vector system [13], [14].

The primary contribution of this paper is examining the use of a single DNN trained for ASR used for both SR and LR tasks. Prior published work has examined the application of DNNs for SR or LR with each method separately. In Section IV we describe initial LR experiments which motivate the focus on two indirect methods. The first indirect method (bottleneck features or BNFs) uses frame-level features extracted from a DNN with a special bottleneck layer [15] and the second indirect method (DNN posteriors) uses posteriors extracted from a DNN to accumulate multi-modal statistics [6]. The features and statistics from these indirect methods are used to train i-vector classifiers for each task and combinations of methods. Contrastive and fusion experiments for well defined SR and LR benchmarks are given in Section V.

II. DNN'S FOR SR AND LR

A DNN classifier is essentially a multi-layer perceptron with more than two hidden layers that typically uses random initialization and stochastic gradient descent to initialize and optimize the weights [1], [16]. For speech applications, the input to a DNN is a stacked set of spectral features (e.g., MFCCs, PLPs) extracted from short (20 ms) segments (frames) of speech. Typically a context of ± 5 to 10 frames around the current input frame are used. The output of the DNN is a prediction of the posterior probability of the target classes for the current input frame (see Fig. 1).

In the direct method for LR and SR, a DNN is used to predict the language or speaker class for a given frame of speech. Since the entire speech waveform is considered to belong to a single class, the frame-level DNN posteriors must be combined to make a single decision score. This can be accomplished either by simply averaging the DNN predictions or by training a secondary classifier, such as a multinomial, that uses statistics

Manuscript received January 25, 2015; revised April 01, 2015; accepted April 02, 2015. Date of publication April 06, 2015; date of current version April 21, 2015. This work was supported by the U. S. Department of Defense under Air Force Contract FA8721-05-C-0002. Opinions, interpretations, conclusions, and recommendations are those of the authors and are not necessarily endorsed by the United States Government. The associate editor coordinating the review of this manuscript and approving it for publication was Prof. Daniel Povey.

F. Richardson and D. Reynolds are with the MIT Lincoln Laboratory, Lexington, MA 02421 USA (e-mail: frichard@ll.mit.edu; dar@ll.mit.edu).

N. Dehak is with the MIT Computer Science and Artificial Intelligence Laboratory, Cambridge, MA 02139 USA (e-mail: najim@csail.mit.edu).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/LSP.2015.2420092

across the whole input derived from the DNN as a single feature vector.

In contrast to the direct method, the indirect method uses a DNN that was trained on a different data set and possibly for a different purpose. In this work, we have used a DNN trained for an ASR task for both LR and SR. The ASR DNN is trained to predict sub-phonetic units or “senones” for each input frame [1]. In the following two subsections we describe how we use the ASR DNN output posteriors and BNFs in the context of an i-vector classifier.

A. DNN Posteriors

A typical i-vector system uses zeroth, first and second order statistics generated using a Gaussian mixture model (GMM) [13] which is commonly referred to as the universal background model (UBM). Statistics are accumulated by first estimating the posterior of each GMM component density for a frame and using these posteriors as weights for accumulating the statistics for each component of the mixture distribution. The zeroth order statistics are the total occupancies across an utterance for each GMM component and the first order statistics are the occupancy weighted accumulations of feature vectors for each component. The i-vector is then computed using a dimension reducing transformation applied to the stacked first order statistics.

An alternate approach to extracting statistics has been proposed in [6]. Statistics are accumulated in the same way as for the GMM but class posteriors from the DNN are used in place of GMM component posteriors. Once the statistics have been accumulated, the i-vector extraction is performed in the same way as it is from the GMM based posteriors. This approach has been shown to give significant gains for both SR and LR [6], [7], [17].

B. DNN Bottleneck Features

A DNN can also be used as a means of extracting features for use by a secondary classifier - including another DNN [18]. This is accomplished by using the activation of one of the DNN’s hidden layers as a feature vector. For some classifiers the dimensionality of the hidden layer is too high and some sort of feature reduction is necessary like LDA or PCA. In [15], a dimension reducing linear transformation is optimized as part of the DNN training by using a special bottleneck hidden layer that has fewer nodes (see Fig. 1). The bottleneck layer uses a linear activation and behaves very much like a LDA or PCA transformation on the activation of the previous layer [15], [19]. Matrix factorization was originally proposed in [19] to reduce the number of parameters of the output layer of the DNN, but in our work we have chosen to use the second to last layer with the hope that the output posterior prediction would not be too adversely affected by the loss of information at the bottleneck layer. BNFs have been shown to work well for both LR [2], [3] and SR [10], [12]. The experiments described in Section V focus on using a single DNN trained for ASR to extract bottleneck features and estimate posteriors. While it may appear that BNFs derived from an ASR trained DNN should have little speaker information, the results below and also reported in [12] indicate that these features still contain speaker-dependent phonetically discriminative information that is beneficial for the SR task.

III. I-VECTOR SYSTEM

In the experiments, an i-vector classifier, configured as described below, was used for baseline and integrated DNN systems. Speech activity segmentation generated using a GMM based speech activity detector was used for all systems. The front-end feature extraction for the baseline LR system uses 7 static cepstra appended to 49 shifted delta cepstra (SDC) for a total of 56 features. The front-end for the baseline SR system uses 20 MFCCs including C0 and their first derivatives for a total of 40 features. Features extracted directly from the DNN bottle neck layer are used for BNF experiments.

All GMM i-vector systems use a 2048 component GMM. For the DNN posterior experiments, the DNN output posteriors are used instead of those from the GMM. In both cases, 600 dimensional i-vectors are extracted from stacked mean vectors which are standardized to $\mathcal{N}(\mathbf{0}, \mathbf{I})$ using diagonal Gaussian parameters μ_{ubm} and Σ_{ubm} . After extraction, i-vectors are length normalized [20]. For both LR and SR, the mean of the enrollment files for each language or speaker are used as the target models. For both tasks, within and across class i-vector covariances (Σ_{wc} and Σ_{ac}) are estimated and used for PLDA scoring. For SR, the i-vectors are whitened using parameters μ_{tot} and Σ_{tot} prior to length normalization.

Hyperparameters for SR and LR tasks (which include μ_{ubm} , Σ_{ubm} , \mathbf{T} , Σ_{wc} , Σ_{ac} , μ_{tot} and Σ_{tot}) were trained using their respective training data with the above i-vector configuration. All LR systems are calibrated using the discriminative Gaussian backend described in [21] which uses both scores and durations for calibration.

IV. INITIAL LR EXPERIMENTS

Our initial DNN experiments focused on the LR task where in contrast to the SR task there are a small number of well defined language classes with a significant amount of data for each one. For these experiments we use a six language sub-set of the NIST 2009 Language Recognition Evaluation (LRE09) corpora that includes Farsi, Hindi, Korean, Mandarin, Russian and Vietnamese [22], [23]. The training partition consists of 20 hours of data per language (10 hours each of VoA and CTS data) for a total of 120 hours of speech. The test data is the subset of the LRE09 evaluation that matches the same six languages. The corpus design was motivated by results reported in Fig. 3 of [4] for roughly the same amount of data per a language. In all the development experiments we calibrate the scores using a cheating backend on the test data and report the C_{avg} at 30, 10 and 3 seconds.

We first focused on the direct approach discussed in Section II. The DNN used here was trained using the training partition of the data set and consisted of 819 input nodes (stack of 21 frames of 13 Gaussianized [24] PLP coefficient and their first and second order derivatives), 2 hidden layers with 2560 nodes per layer, and 6 output nodes. All hidden layers used a sigmoid activation. The DNN training is preformed on an nVidia Tesla K40 GPU using custom software developed at MIT/CSAIL. Scores for each language were generated by averaging the 6 DNN frame-level output log posteriors for each test file.

Results using the direct approach along with the baseline i-vector system are given in Table I. The learning rate for the

TABLE I
INITIAL LR DIRECT DNN PERFORMANCE (C_{avg})

System	30 sec	10 sec	3 sec
Baseline	0.681	2.87	11.4
Direct DNN	2.56	4.24	10.3

TABLE II
INITIAL LR INDIRECT DNN PERFORMANCE (C_{avg})

Features	Posteriors	30 sec	10 sec	3 sec
SDC	GMM	0.684	2.84	11.4
BNF	GMM	0.205	1.24	7.53
SDC	DNN	0.318	1.42	8.75
BNF	DNN	0.304	1.49	8.36

DNN was 0.2 and the frame error rate on a held-out validation set was 56.3% after 10 epochs of training. While there is a substantial degradation in performance relative to the baseline system at the 30 s and 10 s durations using this technique, there is a slight gain at the short 3 s duration for the 2 layer DNN which is consistent with results reported in [4].

We next examined the DNN BNF and the DNN posterior technique discussed in Section II. The DNN for these experiments was trained on 100 hours of Switchboard 1 data [25] using 4,199 state cluster (senone) target labels generated using the Kaldi Switchboard-1 tri4a example system [26]. The same 819 input features were used as the above direct approach DNN. The DNN has 7 hidden layers of 1024 nodes each with the exception of the 6th bottleneck layer which has 64 nodes. All hidden layers use a sigmoid activation function with the exception of 6th layer which is linear [15]. The learning rate for the DNN was 0.2 and the DNN frame error rate on a held-out validation set was 54.1% after 19 epochs of training. This DNN is used for all the experiments reported in Section V.

The results in Table II show that the BNFs together with GMM posteriors give the best performance and that in general all systems using the DNN significantly out perform the baseline system. Based on these results we focused next on applying this ASR DNN to more challenging LR and SR tasks.

During the course of this investigation several other interesting observations were made. Increasing the amount of training data and the number of parameters in the DNN both by up to a factor of three did not significantly improve performance. The inclusion of the bottle neck layer did not adversely affect the DNN posterior based results. Appending delta features to the PLP features gave a small but significant performance gain. Most interestingly, without Gaussianization on the input PLP features LR performance degraded by well over a factor of two (the DNN frame error rate without Gaussianization is 68.9%).

V. LR AND SR BENCHMARK EXPERIMENTS

In this section we present experiments with the indirect DNN approaches on some well defined LR and SR benchmarks. The LR systems were evaluated on the NIST 2011 Language Recognition Evaluation (LRE11) data [27] which covers 24 languages coming from telephone and broadcast audio and has test durations of 3, 10, and 30 seconds. Details on the LR training and

TABLE III
IN-DOMAIN DAC13 RESULTS

Features	Posteriors	EER(%)	DCF*1000
MFCC	GMM	2.71	0.404
MFCC	DNN	2.27	0.336
BNF	GMM	2.00	0.269
BNF	DNN	2.79	0.388

TABLE IV
OUT-OF-DOMAIN DAC13 RESULTS

Features	Posteriors	EER(%)	DCF*1000
MFCC	GMM	6.18	0.642
MFCC	DNN	3.27	0.427
BNF	GMM	2.79	0.342
BNF	DNN	3.97	0.454

TABLE V
LRE11 RESULTS C_{avg}

Features	Posteriors	30s	10s	3s
SDC	GMM	5.26	10.7	20.9
SDC	DNN	4.00	8.21	19.5
BNF	GMM	2.76	6.55	15.9
BNF	DNN	3.79	7.71	18.2
5-way fusion [21]		3.27	6.67	17.1

development data can be found in [21]. The SR systems were trained and evaluated using the 2013 Domain Adaptation Challenge (DAC13) [28]. The DAC13 is a specified set of hyper-parameter, enroll, and test lists developed to exhibit a data domain shift for a SR task and has been reported on in several publications [17], [29], [30]. The same Switchboard trained bottleneck DNN described at the end of Section IV was used in all SR and LR experiments and all senones including those corresponding to non-speech were used.

A. Language Recognition Experiments

The experiments run on the LRE11 task are summarized in Table V with the first row corresponding to the baseline system and the last row corresponding to a fusion of 5 “post-evaluation” systems (see [21] for details). BNFs with GMM posteriors out perform the other systems configurations including the 5 system fusion.

B. Speaker Recognition Experiments

Two sets of experiments were run on the DAC13 corpora: “in-domain” and “out-of-domain”. For both sets of experiments, the UBM and **T** hyper-parameters are trained on Switchboard (SWB) data. The other hyper-parameters (whitening, within, and across covariances) are trained on 2004-2008 speaker recognition evaluation (SRE) data for the in-domain experiments and SWB data for the out-of-domain experiments (see [28] for more details). Tables III and IV summarize the results for the in-domain and out-of-domain experiments with the first row of each table corresponding to the baseline system. While the DNN-posterior technique with MFCCs gives a significant gain over the baseline system for both sets of experiments, as also reported in [6] and [17], an even greater gain is realized using BNF with a GMM. However, using both BNFs and DNN-posteriors degrades performance.

TABLE VI
FUSION OF ALL SYSTEM AND THE TOP 2 SYSTEM ON DAC13.
THE SYSTEM NOTATION USED IS [FEATURE]/[POSTERIOR]

Fusion	Condition	EER(%)	DCF*1000
All 4 systems	In-domain	1.61	0.236
BNF/GMM + MFCC/DNN	In-domain	1.65	0.237
Tandem/GMM	In-domain	1.55	0.229
All 4 systems	Out-of-domain	2.88	0.355
BNF/GMM + MFCC/DNN	Out-of-domain	2.54	0.326
Tandem/GMM	Out-of-domain	2.44	0.323

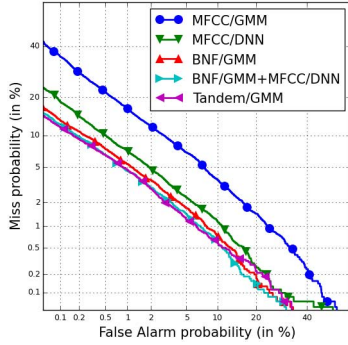


Fig. 2. DAC13 out-of-domain DET plot.

C. Score and Feature Fusion

Scores from the four speaker recognition systems in Tables III and IV were fused by combining them with uniform weights. Out of all possible pair-wise combinations, the BNF/GMM + MFCC/DNN systems yielded the best performance. The results are summarized in Table VI. For the out-of-domain case the 4 system fusion is actually worse than fusing just the BNF/GMM + MFCC/DNN systems perhaps due to the poorer performance of the MFCC/GMM system in this condition. For the in-domain case the BNF/GMM + MFCC/DNN system fusion comes very close to fusing all four systems. While it is possible that better performance could be attained by estimating the optimal weights for combining scores on held-out data or via cross-validation, we believe that the naive fusion using uniform weights is a good indication of how well fusion works between these different systems. The best in-domain score fusion gives a performance gain of almost 20% relative to the BNF/GMM system alone while the best out-of-domain score fusion gives a relative gain of only about 9%.

Also included in Table VI is the result of stacking 20 MFCC features with the 64 BNFs and retraining the GMM I-vector system with the resulting 84 tandem features [31]. The performance for the tandem feature system is slightly better than score fusion for the DAC13 task. The DAC13 out-of-domain DET plots are shown in Fig. 2.

Score fusion experiments using the four language recognition systems in Table V were carried out by training a discriminative backend on the development data over all two system combinations and comparing the top performing pair to the fusion of all four systems. The results are summarized in Table VII. As in the DAC13 fusion experiments, the BNF/GMM + SDC/DNN gave the best performance of all two system combinations.

TABLE VII
LRE11 FUSION C_{avg}

Fusion	30s	10s	3s
All 4 systems	2.22	5.41	14.5
BNF/GMM + SDC/DNN	2.31	5.69	14.7
Tandem/GMM	2.67	6.71	15.9

While the fusion gains are relatively modest (roughly a 10% relative improvement across the durations), the fusion of just the BNF/GMM + SDC/DNN is only slightly worse than the fusion of all four systems.

The tandem system performs worse than score fusion on the LRE11 task but is on par with the BNF/GMM system. This may be because the features used for score fusion and for the tandem features are not the same. The MFCC features used in the tandem system perform well on the SR task but are not as suited to the LR task as the SDC features used in the SDC/DNN system. However, the tandem/GMM system's result suggests that one could use the same tandem feature representation for both LR and SR and still realize a gain on the SR task. This may be of interest in situations where i-vectors are extracted with one set of hyper parameters and then used for both the LR and SR tasks. Estimating μ_{ubm} , Σ_{ubm} and T parameters that yield good performance on both tasks is an area for future research.

VI. CONCLUSIONS

This paper has described the development of a DNN BNF i-vector system and demonstrated substantial performance gains when applying the system to both the DAC13 SR and LRE11 LR benchmarks. For the DAC13 task the BNF/GMM system was shown to reduce the error rates of the baseline MFCC/GMM system by 26% for EER and 33% for DCF for the in-domain task and 55% for EER and 47% for DCF for the out-of-domain task. On LRE11, the same BNFs decreased C_{avg} at 30 s, 10 s, and 3 s durations by 48%, 39%, and 24%, respectively, and even out performed a 5 system fusion of acoustic and phonetic based recognizers [21].

Further reductions in error were demonstrated on the DAC13 SR task using score fusion or tandem features. Fusing the BNF/GMM and MFCC/DNN system scores reduces the error rates relative to the BNF/GMM system by 18% for EER and 12% for DCF for the in-domain task and by 9% for EER and 5% for DCF for the out-of-domain task. Using tandem features lead to a larger reduction in error rate of 23% for EER and 15% for DCF for the in-domain task and 13% for EER and 6% for DCF by for the out-of-domain task. Score fusion on the LRE11 task lead to 16%, 13% and 8% reduction in C_{avg} on the 30 s, 10 s and 3 s durations conditions. While the tandem features did not lead to significant changes in performance on the LRE11 task, their good performance on DAC13 suggests the possibility of a single tandem front-end and a single I-vector extractor for both LR and SR applications.

ACKNOWLEDGMENT

The authors would like to thank Patrick Cardinal, Yu Zhang and Ekapol Chuangsuwanich at MIT CSAIL for sharing their DNN expertise and GPU optimized DNN training software.

REFERENCES

- [1] G. Hinton *et al.*, “Deep neural networks for acoustic modeling in speech recognition,” *IEEE Signal Process. Mag.*, pp. 82–97, Nov. 2012.
- [2] Y. Song, B. Jiang, Y. Bao, S. Wei, and L.-R. Dai, “I-vector representation based on bottleneck features for language identification,” *Electron. Lett.*, pp. 1569–1580, 2013.
- [3] P. Matejka, L. Zhang, T. Ng, H. S. Mallidi, O. Glembek, J. Ma, and B. Zhang, “Neural network bottleneck features for language identification,” in *Proc. IEEE Odyssey*, 2014, pp. 299–304.
- [4] I. Lopez-Moreno, J. Gonzalez-Dominguez, O. Plchot, D. Martinez, J. Gonzalez-Rodriguez, and P. Moreno, “Automatic language identification using deep neural networks,” in *Proc. ICASSP*, 2014, pp. 5374–5378.
- [5] T. Yamada, L. Wang, and A. Kai, “Improvement of distant-talking speaker identification using bottleneck features of dnn,” in *Proc. Interspeech*, 2013, pp. 3661–3664.
- [6] Y. Lei, N. Scheffer, L. Ferrer, and M. McLaren, “A novel scheme for speaker recognition using a phonetically-aware deep neural network,” in *Proc. ICASSP*, 2014, pp. 1714–1718.
- [7] Y. Lei, L. Ferrer, A. Lawson, M. McLaren, and N. Scheffer, “Application of convolutional neural networks to language identification in noisy conditions,” in *Proc. IEEE Odyssey*, 2014, pp. 287–292.
- [8] P. Kenny, V. Gupta, T. Stafylakis, P. Ouellet, and J. Alam, “Deep neural networks for extracting baum-welch statistics for speaker recognition,” in *Proc. IEEE Odyssey*, 2014, pp. 293–298.
- [9] O. Ghahabi and J. Hernando, “I-vector modeling with deep belief networks for multi-session speaker recognition,” in *Proc. IEEE Odyssey*, 2014, pp. 305–310.
- [10] S. Yaman, J. Pelecanos, and R. Sarikaya, “Bottleneck features for speaker recognition,” in *Proc. IEEE Odyssey*, 2012.
- [11] E. Variani, X. Lei, E. McDermott, I. Lopez-Moreno, and J. Gonzalez-Dominguez, “Deep neural networks for small footprint text-dependent speaker verification,” in *Proc. ICASSP*, 2014, pp. 4080–4084.
- [12] A. K. Sarkar, C.-T. Do, V.-B. Le, and C. Barras, “Combination of cepstral and phonetically discriminative features for speaker verification,” *IEEE Signal Process. Lett.*, vol. 21, no. 9, pp. 1040–1044, Sep. 2014.
- [13] N. Dehak, P. Kenny, R. Dehak, P. Ouellet, and P. Dumouchel, “Front end factor analysis for speaker verification,” *IEEE Trans. Audio, Speech, Lang. Process.*, vol. 19, no. 4, pp. 788–798, May 2011.
- [14] N. Dehak, P. Torres-Carrasquillo, D. Reynolds, and R. Dehak, “Language recognition via ivectors and dimensionality reduction,” in *Proc. Interspeech*, 2011, pp. 857–860.
- [15] Y. Zhang, E. Chuangsuwanich, and J. Glass, “Extracting deep neural network bottleneck features using low-rank matrix factorization,” in *Proc. ICASSP*, 2014, pp. 185–189.
- [16] L. Deng, G. Hinton, and B. Kingsbury, “New types of deep neural network learning for speech recognition and related applications: An overview,” in *Proc. ICASSP*, 2013.
- [17] D. Garcia-Romero, X. Zhang, A. McCree, and D. Povey, “Improving speaker recognition performance in the domain adaptation challenge using deep neural networks,” in *Proc. IEEE SLT Workshop*, 2014.
- [18] K. Vesely, M. Karafiat, and F. Grezl, “Convolutional bottleneck network features for lvsr,” in *Proc. IEEE ASRU*, 2011, pp. 42–47.
- [19] T. N. Sainath, B. Kingsbury, V. Sindhwani, E. Arisoy, and B. Ramabhadran, “Low-rank matrix factorization for deep neural network training with high-dimensional output targets,” in *Proc. ICASSP*, 2013.
- [20] D. Garcia-Romero and C. Y. Espy-Wilson, “Analysis of i-vector length normalization in speaker recognition systems,” in *Proc. Interspeech*, 2011, pp. 249–252.
- [21] E. Singer, P. Torres-Carrasquillo, D. Reynolds, A. McCree, F. Richardson, N. Dehak, and D. Sturim, “The MITLL NIST LRE 2011 language recognition system,” in *Proc. IEEE Odyssey*, 2011, pp. 209–215.
- [22] “The 2009nist language recognition evaluation plan,” 2009 [Online]. Available: <http://www.itl.nist.gov/iad/mig/tests/lre/2009/>
- [23] P. Torres-Carrasquillo, E. Singer, T. Gleason, A. McCree, D. A. Reynolds, F. Richardson, and D. Sturim, “The MITLL NIST LRE 2009 language recognition system,” in *Proc. ICASSP*, 2010, pp. 4994–4997.
- [24] B. Xiang, U. Chaudhari, J. Navratil, G. Ramaswamy, and R. Gopinath, “Short-time Gaussianization for robust speaker verification,” in *Proc. ICASSP*, 2002.
- [25] J. Godfrey, E. Holliman, and J. McDaniel, “Switchboard: Telephone speech corpus for research and development,” in *Proc. ICASSP*, 1992, pp. 517–520.
- [26] D. Povey *et al.*, “The kald speech recognition toolkit,” in *Proc. IEEE ASRU*, 2011.
- [27] “The 2011 nist language recognition evaluation plan,” 2011 [Online]. Available: <http://www.nist.gov/itl/iad/mig/lre11.cfm>
- [28] S. H. Shum, D. A. Reynolds, D. Garcia-Romero, and A. McCree, “Unsupervised clustering approaches for domain adaptation in speaker recognition systems,” in *Proc. IEEE Odyssey*, 2014, pp. 265–272.
- [29] H. Aronowitz, “Inter dataset variability compensation for speaker recognition,” in *Proc. ICASSP*, 2014.
- [30] O. Glembek, J. Ma, P. Matejka, B. Zhang, O. Plchot, L. Burget, and S. Matsoukas, “Domain adaptation via within-class covariance correction in i-vector based speaker recognition systems,” in *Proc. ICASSP*, 2014.
- [31] H. Hermansky, D. P. W. Ellis, and S. Sharma, “Tandem connectionist feature extraction for conventional HMM systems,” in *Proc. ICASSP*, 2000.